

## 04. Data Access and Manipulation

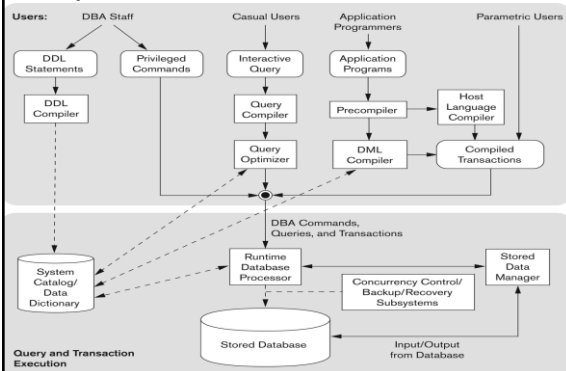
1

### Introduction: What is SQL

- Expanded: Structured Query Language
  - Initially called SEQUEL (Structured English Query Language)
- A special purpose programming language for managing database management systems
- Both a declarative and procedural language
- SQL considered a major reason for the success of relational databases:
  - It became a standard for relational databases => fewer migration worries
  - Users can write statements in a database application program that accesses two or more databases without having to alter language
- There still are differences in various dialects of SQL ... but if users stick to the 'Standard SQL' and relational systems faithfully support it

2

### Component Modules of a DBMS



4

### Data Definition and Data Types

- Keywords: CREATE, DROP, and ALTER the descriptions of the tables (relations) of a database
- Required: descriptor of structure of table to be created (Data Types)
- Typically the role of the DBA and any other user explicitly authorized by DBA

4

### Standard SQL Datatypes

- Character String:
  - CHAR(N), CHARACTER(N)
  - VARCHAR(n)
- Numeric:
  - INTEGER (OR INT), SMALLINT,
  - FLOAT (OR REAL), DOUBLE PRECISION
  - DECIMAL(I,J)
- Bit-string
  - BIT(n) – for fixed length
  - BIT VARYING(n) – n is the max
- BOOLEAN – values True, False and Unknown (in cases of 3-valued logic)
- DATE – handles values of format YYYY-MM-DD
- TIME – HH:MM:SS
- TIMESTAMP – Includes both date and time

5

### CREATE SCHEMA

- Specifies a new database schema by giving it a name
- Database Schema = a set of relation schemas

6

## CREATE TABLE

- Specifies a new relation by giving it a name, and specifying each of its attributes and their data types (INTEGER, FLOAT, DECIMAL(i,j), CHAR(n), VARCHAR(n))
- A constraint NOT NULL may be specified on an attribute

```
CREATE TABLE DEPARTMENT
(DNAME VARCHAR(10) NOT NULL,
DNUMBER INTEGER NOT NULL,
MGRSSN CHAR(9),
MGRSTARTDATE Date);
```

7

## CREATE TABLE- Additional

- CREATE TABLE command for specifying the primary key attributes, secondary keys, and referential integrity constraints (foreign keys).
- Key attributes can be specified via the PRIMARY KEY and UNIQUE phrases

```
CREATE TABLE DEPT
( DNAME VARCHAR(10) NOT NULL,
  DNUMBER INTEGER NOT NULL,
  MGRSSN CHAR(9),
  MGRSTARTDATE Date,
  PRIMARY KEY (DNUMBER),
  UNIQUE (DNAME),
  FOREIGN KEY (MGRSSN) REFERENCES EMP(SSN) );
```

8

## DROP TABLE

- Used to remove a relation (base table) *and its definition*
- The relation can no longer be used in queries, updates, or any other commands since its description no longer exists
- Example:

**DROP TABLE DEPENDENT;**

9

## ALTER TABLE

- Used to add an attribute to one of the base relations
- The new attribute will have NULLs in all the tuples of the relation right after the command is executed; hence, the NOT NULL constraint is *not allowed* for such an attribute
- Examples:
  - ALTER TABLE Employee ADD Job VARCHAR(12);
  - ALTER TABLE Employee DROP Job
- The database users must still enter a value for the new attribute JOB for each EMPLOYEE tuple. This can be done using the UPDATE command.

10

## REFERENTIAL INTEGRITY OPTIONS

- We can specify RESTRICT, CASCADE, SET NULL or SET DEFAULT on referential integrity constraints (foreign keys)

```
CREATE TABLE DEPT
( DNAME VARCHAR(10) NOT NULL,
  DNUMBER INTEGER NOT NULL,
  MGRSSN CHAR(9),
  MGRSTARTDATE CHAR(9),
  PRIMARY KEY (DNUMBER),
  UNIQUE (DNAME),
  FOREIGN KEY (MGRSSN) REFERENCES EMP
ON DELETE SET DEFAULT ON UPDATE CASCADE );
```

11

## Retrieval Queries in SQL

- SQL has one basic statement for retrieving information from a database; the SELECT statement
- Important distinction between SQL and the formal relational model; SQL allows a table (relation) to have two or more tuples that are identical in all their attribute values
- SQL relations can be constrained to be sets by specifying PRIMARY KEY or UNIQUE attributes, or by using the DISTINCT option in a query

12

## Retrieval Queries in SQL

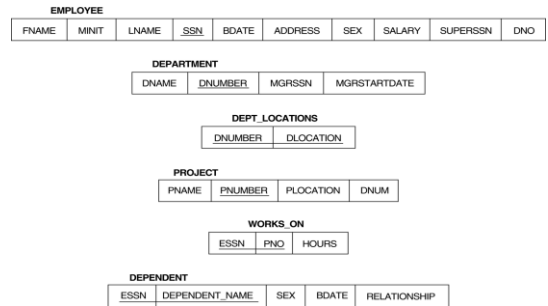
- Basic form of the SQL SELECT statement is called a *mapping* or a *SELECT-FROM-WHERE block*

**SELECT** <attribute list>  
**FROM** <table list>  
**WHERE** <condition>

- <attribute list> is a list of attribute names whose values are to be retrieved by the query
- <table list> is a list of the relation names required to process the query
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query

13

## Relational Database Schema



14

## Specifying Updates in SQL

- There are three SQL commands to modify the database state; INSERT, DELETE, and UPDATE

15

## INSERT

- In its simplest form, it is used to add one or more tuples to a relation
- Attribute values should be listed in the same order as the attributes were specified in the CREATE TABLE command

16

## INSERT (cont.)

- Example U1:

**INSERT INTO Employee**

**VALUES ('Richard', 'K', 'Waweru', '653298653', '12 Dec 1962', 'P O Box 98 Thika', 'M', 37000, '987654321', '4');**

17

## Alternative INSERT

- Specifies explicitly the attribute names that correspond to the values in the new tuple
- Attributes with NULL values can be left out
- Example: Insert a tuple for a new EMPLOYEE for whom we only know the FNAME, LNAME, and SSN attributes.

U1a:

**INSERT INTO Employee (FName, LName, DNo, SSN )  
 VALUES ('Patrick', 'Mwachiro', '4', '653298654');**

18

## INSERT

- **Important Note:** Only the constraints specified in the DDL commands are automatically enforced by the DBMS when updates are applied to the database
- Another variation of INSERT allows insertion of *multiple tuples* resulting from a query into a relation
- **Example:** Suppose there's need to create a temporary table that has the name, number of employees, and total salaries for each department. A table DEPTS\_INFO is created by U3A, and is loaded with the summary information retrieved from the database by the query in U3B.

19

## INSERT

### U3a:

```
CREATE TABLE Depts_Info  
(Dept_Name VARCHAR(15), No_Of_Emps Integer,  
  Total_Sal Integer);
```

### U3b:

```
INSERT INTO Depts_Info ( Dept_Name, No_Of_Emps,  
  Total_Sal )  
SELECT DNAME, Count(*), Sum(Employee.Salary)  
FROM Department, Employee  
WHERE DNUMBER =DNo  
GROUP BY DNAME;
```

20

## DELETE

- Removes tuples from a relation
- Includes a WHERE-clause to select the tuples to be deleted
- Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint)
- A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table
- The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause
- Referential integrity should be enforced

21

## Delete

### U4a:

```
DELETE *  
FROM Employee  
WHERE LName='Waweru';
```

### U4b:

```
DELETE *  
FROM Emp1  
WHERE SSN='123456789';
```

### U4c:

```
DELETE *  
FROM Emp3  
WHERE DNo IN  
(Select DNumber From Department Where DName='Research');
```

### U4d:

```
DELETE *  
FROM Emp4;
```

22

## UPDATE

- Used to modify attribute values of one or more selected tuples
- A WHERE-clause selects the tuples to be modified
- An additional SET-clause specifies the attributes to be modified and their new values
- Each command modifies tuples *in the same relation*
- Referential integrity should be enforced<sup>23</sup>

23

## UPDATE (cont.)

- **Example:** Change the location and controlling department number of project number 10 to 'Garissa' and 5, respectively.

### U5:

```
UPDATE Proj SET PLocation = 'Garissa', Dnum  
= 5  
WHERE PNumber=10;
```

24

## UPDATE

**Example (U6):** Give all employees in the 'Research' department a 10% raise in salary.

**UPDATE Emp2 SET Salary = Salary\*1.1  
WHERE DNO IN  
(SELECT DNumber  
FROM DEPARTMENT  
WHERE DName = 'Research');**

- In this request, the modified SALARY value depends on the original SALARY value in each tuple

25

## Populated Database

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
John	B		Smith	123456789	1965-01-09	731 Funderd, Houston, TX	M	30000	333445555	5
Franklin	T		Wong	333445555	1965-10-08	438 Voss, Houston, TX	M	49000	333445555	5
Allen	J	Wayne	999997777	1965-07-19	3321 Center, Dallas, TX	F	25000	997554321	4	
James	B	Wayne	997554321	1941-06-25	251 River, Dallas, TX	F	49000	999997777	4	
Robert	A	Wayne	999994444	1962-09-15	375 Pine, San, Houston, TX	M	38000	333445555	5	
James	A	Wayne	333445555	1972-07-15	567 Pine, Houston, TX	F	25000	333445555	5	
James	E	Wayne	997554321	1965-05-20	480 Llama, Houston, TX	M	25000	997554321	4	
James	E	Wayne	999997777	1937-11-10	480 Llama, Houston, TX	M	25000	null	1	

DEPT LOCATIONS			
DEPARTMENT	DNAME	DNUMBER	LOCATION
Administration	5	997554321	1
Research	1	999997777	5

WORKS ON	ESSN	PNO	HOURS
123456789	1	10.5	
123456789	2	7.5	
333445555	3	20.0	
455432109	1	20.0	
455432109	2	20.0	
333445555	3	10.0	
333445555	4	10.0	
333445555	10	10.0	
333445555	20	10.0	
999997777	30	30.0	
999997777	10	10.0	
997554321	10	5.0	
997554321	20	10.0	
997554321	30	10.0	
999997777	20	10.0	

PROJECT	FNAME	PNUMBER	PLOCATION	PRJLM
ProjectX	1	Dallas	5	
ProjectY	2	San Antonio	5	
ProjectZ	3	Houston	5	
Compensation	10	Dallas	4	
Reorganization	20	Houston	1	
Reorganization	30	Dallas	4	

DEPENDENT	ESSN	DEPENDENT NAME	SEX	BDATE	RELATIONSHIP
333445555	Allen	Franklin	M	1965-10-08	DAUGHTER
333445555	Franklin	James	M	1965-10-08	SON
333445555	James	James	F	1965-10-08	DAUGHTER
997554321	Allen	James	M	1941-06-25	SON
123456789	Allen	James	M	1965-01-09	SON
123456789	Allen	James	F	1965-01-09	DAUGHTER

## Simple SQL Queries

- Basic SQL queries correspond to using the SELECT, PROJECT, and JOIN operations of the relational algebra
- All subsequent examples use the COMPANY database
- Example of a simple query on *one* relation
- Query 0:** Retrieve the birthdate and address of the employee whose name is 'John'.

**Q0:SELECT BDATE, ADDRESS  
FROM EMPLOYEE  
WHERE FNAME='John'**

27

## Simple SQL Queries

- Query 1:** Retrieve the name and address of all employees who work for the 'Research' department.

**Q1**

**SELECT FName, LName, Address  
FROM Employee, Department  
WHERE DName='Research' And DNumber=DNo;**

- Similar to a SELECT-PROJECT-JOIN sequence of relational algebra operations
- (DNAME='Research') is a *selection condition* (corresponds to a SELECT operation in relational algebra)
- (DNUMBER=DNO) is a *join condition* (corresponds to a JOIN operation in relational algebra)

28

## Simple SQL Queries (cont.)

- Query 2:** For every project located in 'Naivasha', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

**SELECT PNumber, DNum, LName, BDate, Address  
FROM Project, Department, Employee  
WHERE DNum=DNumber AND MgrSSN=SSN AND PLOCATION='Naivasha';**

- The join condition **DNUM=DNUMBER** relates a project to its controlling department
- The join condition **MGRSSN=SSN** relates the controlling department to the employee who manages that department

29

## Aliases, \* and DISTINCT, Empty WHERE-clause

- In SQL, we can use the same name for two (or more) attributes as long as the attributes are in *different relations*
- A query that refers to two or more attributes with the same name must *qualify* the attribute name with the relation name by *prefixing* the relation name to the attribute name

**Example:**

- EMPLOYEE.NAME, DEPARTMENT.NAME**

30

## ALIASES

- Some queries need to refer to the same relation twice
- In this case, *aliases* are given to the relation name
- **Query 8:** For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

**Q8:**

```
SELECT E.Fname, E.LName, S.FName, S.Lname
FROM Employee E, Employee S
WHERE E.SuperSSN=S.SSN;
```

- The alternate relation names E and S are called *aliases* for the EMPLOYEE relation
- We can think of E and S as two *different copies* of EMPLOYEE; E represents employees in role of *supervisees* and S represents employees in role of *supervisors*

31

## ALIASES (cont.)

Aliasing can also be used in any SQL query for convenience. Can also use the AS keyword to specify aliases

```
SELECT E.Fname, E.LName, S.FName, S.Lname
FROM Employee AS E, Employee AS S
WHERE E.SuperSSN=S.SSN;
```

32

## UNSPECIFIED WHERE-clause

- A *missing WHERE-clause* indicates no condition; hence, *all tuples* of the relations in the FROM-clause are selected
- This is equivalent to the condition WHERE TRUE
- **Query 9:** Retrieve the SSN values for all employees.

**Q9:**

```
SELECT SSN
FROM EMPLOYEE
```

- If more than one relation is specified in the FROM-clause *and* there is no join condition, then the *CARTESIAN PRODUCT* of tuples is selected

**Q10:**

```
SELECT SSN, DName
FROM Employee, Department;
```

33

## USE OF \*

- To retrieve all the attribute values of the selected tuples, a \* is used, which stands for *all the attributes*
- Examples:

**Q1C:**

```
SELECT *
FROM EMPLOYEE
WHERE DNO=5
```

**Q1D:**

```
SELECT *
FROM EMPLOYEE, DEPARTMENT
WHERE DNAME='Research' AND
DNO=DNUMBER
```

34

## USE OF DISTINCT

- SQL does not treat a relation as a set; *duplicate tuples can appear*
- To eliminate duplicate tuples in a query result, the keyword **DISTINCT** is used
- For example, the result of Q11 may have duplicate SALARY values whereas Q11A does not have any duplicate values

**Q11:**

```
SELECT SALARY FROM EMPLOYEE
```

**Q11A:**

```
SELECT DISTINCT SALARY FROM EMPLOYEE
```

35

## SET OPERATIONS

- SQL has directly incorporated some set operations
- Operations:
  - Union
  - Minus - Difference
  - Intersect
- The resulting relations of these set operations are sets of tuples; *duplicate tuples are eliminated from the result*
- The set operations apply only to *union compatible relations*; the two relations must have the same attributes and the attributes must appear in the same order

36

## SET OPERATIONS

- **Query 4:** Make a list of all project numbers for projects that involve an employee whose last name is 'Omondi' as a worker or as a manager of the department that controls the project.

**Q4:**

```
(SELECT PName
FROM Project, Department, Employee
WHERE DNUM=DNumber AND MGRSSN=SSN AND LNAME='Omondi')
UNION
(SELECT PName
FROM Project, Works_On, Employee
WHERE PNumber=PNo AND ESSN=SSN AND LName='Omondi');)
```

37

## NESTING OF QUERIES

- A complete SELECT query, called a *nested query*, can be specified within the WHERE-clause of another query, called the *outer query*
- Many of the previous queries can be specified in an alternative form using nesting
- **Query 1:** Retrieve the name and address of all employees who work for the 'Research' department.

**Q1 - Nested**

```
SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE
WHERE DNO IN
      (SELECT DNUMBER
       FROM DEPARTMENT
       WHERE DNAME='Research')
```

38

## NESTING OF QUERIES

- The nested query selects the number of the 'Research' department
- The outer query select an EMPLOYEE tuple if its DNO value is in the result of either nested query
- In general, we can have several levels of nested queries

39

## EXPLICIT SETS

- It is also possible to use an **explicit (enumerated) set of values** in the WHERE-clause rather than a nested query
- **Query 13:** Retrieve the social security numbers of all employees who work on project number 1, 2, or 3.

```
Q13:  SELECT  DISTINCT ESSN
      FROM    WORKS_ON
      WHERE   PNO IN (1, 2, 3)
```

40

## SUBSTRING COMPARISON

- The **LIKE** comparison operator is used to compare partial strings
- Two reserved characters are used: '%' (or '\*' in some implementations) replaces an arbitrary number of characters, and '\_' replaces a single arbitrary character

41

## SUBSTRING COMPARISON

- **Query 12:** Retrieve all employees whose address is in Nairobi. Here, the value of the ADDRESS attribute must contain the substring 'Nairobi'.

**Q12:**

```
SELECT FName, LName, Address
FROM Employee
WHERE Address LIKE '*Nairobi*';
```

42

## ARITHMETIC OPERATIONS

- The standard arithmetic operators '+', '-', '\*', and '/' (for addition, subtraction, multiplication, and division, respectively) can be applied to numeric values in an SQL query result
- Query 13: Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

Q13:

```
SELECT FNAME, LNAME, 1.1*SALARY
FROM EMPLOYEE, WORKS_ON, PROJECT
WHERE SSN=ESSN AND PNO=PNUMBER AND
PNAME='ProductX'
```

43

## Ordering Query Results

- The **ORDER BY** clause is used to sort the tuples in a query result based on the values of some attribute(s)
- Query 15: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

Q15:

```
SELECT DName, LName, FName, PName
FROM Department, Employee, Works_On, Project
WHERE DNumber=DNo AND SSN=ESSN AND
PNO=PNUMBER
ORDER BY DName, LName, FName;
```

44

## Ordering Query Results

- The default order is in ascending order of values
- To specify use the keywords:
  - **DESC** for descending order;
  - **ASC** for ascending order, though it is the default

45

## NULLS IN SQL QUERIES

- SQL allows queries that check if a value is NULL (missing or undefined or not applicable)
- SQL uses **IS** or **IS NOT** to compare NULLs because it considers each NULL value distinct from other NULL values, so equality comparison is not appropriate.
- Query 14: Retrieve the names of all employees who do not have supervisors.

Q14:

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE SUPERSSN IS NULL
```

Note: If a join condition is specified, tuples with NULL values for the join attributes are not included in the result

46

## AGGREGATE FUNCTIONS

- Include **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- Query: Find the sum of salaries, maximum salary, the minimum salary, and the average salary among all employees.

Q19:

```
SELECT SUM(Salary), MAX(Salary), MIN(Salary), AVG(Salary)
FROM EMPLOYEE
```

47

## AGGREGATE FUNCTIONS (cont.)

- Query 20: Find the maximum salary, the minimum salary, and the average salary among employees who work for the 'Research' department.

Q20:

```
SELECT MAX(SALARY), MIN(SALARY), AVG(SALARY)
FROM EMPLOYEE, DEPARTMENT
WHERE DNO=DNUMBER AND DNAME='Research'
```

48



## AGGREGATE FUNCTIONS (cont.)

- Queries 21 and 22: Retrieve the total number of employees in the company (Q21), and the number of employees in the 'Research' department (Q22).

Q21:

```
SELECT COUNT (*) FROM EMPLOYEE
```

Q22:

```
SELECT COUNT (*)  
FROM EMPLOYEE, DEPARTMENT  
WHERE DNO=DNUMBER AND DNAME='Research'
```

49

## GROUPING

- In many cases, we want to apply the aggregate functions *to subgroups of tuples in a relation*
- Each subgroup of tuples consists of the set of tuples that have *the same value* for the *grouping attribute(s)*
- The function is applied to each subgroup independently
- SQL has a **GROUP BY**-clause for specifying the grouping attributes, which *must also appear in the SELECT-clause*

## GROUPING

- Query 24: For each department, retrieve the department number, the number of employees in the department, and their average salary.

Q24:

```
SELECT DNO, COUNT (*), AVG (SALARY)  
FROM EMPLOYEE  
GROUP BY DNO
```

- In Q20, the EMPLOYEE tuples are divided into groups – each group having the same value for the grouping attribute DNO
- The COUNT and AVG functions are applied to each such group of tuples separately
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples
- A join condition can be used in conjunction with grouping <sup>51</sup>

## GROUPING

- Query 25: For each project, retrieve the project number, project name, and the number of employees who work on that project.

Q25:

```
SELECT PNumber, PName, Count(*) AS [Number of  
Workers]  
FROM Project, Works_On  
WHERE PNumber=PNO  
GROUP BY PNumber, PName
```

- In this case, the grouping and functions are applied *after* the joining of the two relations <sup>52</sup>

## Summary of SQL Queries

- A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory. The clauses are specified in the following order:

```
SELECT    <attribute list>  
FROM      <table list>  
[WHERE    <condition>]  
[GROUP BY <grouping attribute(s)>]  
[HAVING   <group condition>]  
[ORDER BY <attribute list>]
```

53

## Summary of SQL Queries

- The SELECT-clause lists the attributes or functions to be retrieved
- The FROM-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The WHERE-clause specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- GROUP BY specifies grouping attributes
- HAVING specifies a condition for selection of groups
- ORDER BY specifies an order for displaying the result of a query
- A query is evaluated by first applying the WHERE-clause, then GROUP BY and HAVING, and finally the SELECT-clause

54